

```

# Load necessary libraries
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo

library(quantmod)

## Loading required package: xts
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
## Loading required package: TTR
library(ggplot2)
library(tseries)
# Define the coefficients for the MA(2) process
theta1 <- -2
theta2 <- 1.35

# Theoretical Autocorrelation for an MA(2) process
rho <- c(1,
        (theta1 + theta1 * theta2) / (1 + theta1^2 + theta2^2),
        (theta2) / (1 + theta1^2 + theta2^2),
        rep(0, 8)) # Autocorrelations for lags > 2 are 0

# Print the theoretical autocorrelation
print(rho)

## [1] 1.0000000 -0.6888970 0.1978747 0.0000000 0.0000000 0.0000000
## [7] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000

# Set seed for reproducibility
set.seed(0)

# Simulate the MA(2) process
t_max <- 100
epsilon <- rnorm(t_max + 2) # Generate white noise terms
y <- rep(0, t_max)

# Generate the time series according to the MA(2) formula
for (t in 3:(t_max + 2)) {
  y[t - 2] <- 0.7 - 2 * epsilon[t - 1] + 1.35 * epsilon[t - 2] + epsilon[t]
}

# Compute the sample autocorrelation function
sample_acf <- acf(y, lag.max=10, plot=FALSE)


# Print the sample autocorrelation function
print(sample_acf$acf)

```

```
## , , 1
##
##      [,1]
## [1,] 1.00000000
## [2,] -0.67851999
## [3,] 0.26881020
## [4,] -0.16730077
## [5,] 0.08060558
## [6,] 0.09770112
## [7,] -0.17831442
## [8,] 0.15498987
## [9,] -0.12203414
## [10,] 0.14946893
## [11,] -0.18055663

# You can also plot the ACF for visual analysis
acf(y, lag.max=10, main="Sample Autocorrelation Function")
```

Sample Autocorrelation Function



To compare the theoretical and sample ACF visually, you can plot them together
Note: You would need to adjust the indices and scaling to make the comparison.

#5

Assuming y is the time series data from the previous exercise

a. Estimate an MA(2) process

```

fit <- Arima(y, order=c(0,0,2))
summary(fit)

## Series: y
## ARIMA(0,0,2) with non-zero mean
##
## Coefficients:
##          ma1      ma2      mean
##       -1.3615  0.5373  0.6979
## s.e.    0.0843  0.0989  0.0207
##
## sigma^2 = 1.374: log likelihood = -157.38
## AIC=322.77   AICc=323.19   BIC=333.19
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.02093061 1.154666 0.9369927 -35.0704 248.9708 0.3156087
##              ACF1
## Training set 0.01791618

# b. Compute the forecasts
# Since we know the last two errors, we can use them for forecasting
# The 1-step ahead forecast will use both errors
# The 2-step and 3-step ahead forecasts will be the mean of the process for an MA(2)

last_error <- c(0.4, -1.2) # Known last errors
forecasts <- numeric(3)
forecasts[1] <- fit$coef[1] * last_error[1] + fit$coef[2] * last_error[2] + mean(fit$residuals)
forecasts[2] <- mean(fit$residuals) # 2-step ahead forecast is the mean for MA(2)
forecasts[3] <- mean(fit$residuals) # 3-step ahead forecast is the same

# Print the forecasts
forecasts

## [1] -1.16843074  0.02093061  0.02093061

#6

# Assuming that you have a time series 'y' that follows the MA(1) process with theta = 0.8
# First, simulate an MA(1) process with theta = 0.8
set.seed(123) # for reproducibility
epsilon <- rnorm(1000) # white noise
y <- filter(epsilon, sides=1, filter=c(0.8,1)) # MA(1) process

# Fit the MA(1) model to the simulated data
ma_fit <- Arima(y, order=c(0,0,1), include.mean=TRUE)

# Print the summary of the fitted model
summary(ma_fit)

## Series: y
## ARIMA(0,0,1) with non-zero mean
##
## Coefficients:

```

```
##          ma1      mean
##          0.7912  0.0297
## s.e.    0.0206  0.0561
##
## sigma^2 = 0.9843: log likelihood = -1409.11
## AIC=2824.23   AICc=2824.25   BIC=2838.95
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0002278621 0.9911321 0.7955465 61.86523 197.7886 0.763285
##              ACF1
## Training set -0.01816309


# Forecast the next 10 values
forecasts <- forecast(ma_fit, h=10)

# Print the forecasts
print(forecasts)

##          Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## 1001      0.34250539 -0.9289549 1.613966 -1.602025 2.287036
## 1002      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1003      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1004      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1005      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1006      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1007      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1008      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1009      0.02967106 -1.5916254 1.650967 -2.449888 2.509230
## 1010      0.02967106 -1.5916254 1.650967 -2.449888 2.509230

# Plot the forecasts
plot(forecasts)
```

Forecasts from ARIMA(0,0,1) with non-zero mean



```
#10
```

```
# Download stock prices
```

```
getSymbols("AAPL", from = "2020-01-01", to = "2024-01-01")
```


```
## [1] "AAPL"
```

```
aapl_returns <- dailyReturn(AAPL)
```

```
# Obtain autocorrelation function
```

```
aapl_acf <- acf(aapl_returns, lag.max = 20, plot = TRUE)
```

Series aapl_returns



```
# Fit ARIMA model
aapl_fit <- auto.arima(aapl_returns)

# Print summary
summary(aapl_fit)


## Series: aapl_returns
## ARIMA(0,0,1) with non-zero mean
##
## Coefficients:
##          ma1      mean
##        -0.1308  0.0012
## s.e.    0.0308  0.0006
##
## sigma^2 = 0.0004396:  log likelihood = 2461.57
## AIC=-4917.14   AICc=-4917.12   BIC=-4902.4
##
## Training set error measures:
##              ME          RMSE          MAE MPE MAPE          MASE          ACF1
## Training set 1.045879e-06 0.02094537 0.01494225 NaN  Inf  0.6710695 -0.002330375

# Forecast the next 10 periods
aapl_forecasts <- forecast(aapl_fit, h = 10)

# Print the forecasts
print(aapl_forecasts)
```

| ## | Point Forecast | Lo 80 | Hi 80 | Lo 95 | Hi 95 |
|---------|----------------|-------------|------------|-------------|------------|
| ## 1007 | 0.002020807 | -0.02484849 | 0.02889010 | -0.03907223 | 0.04311385 |
| ## 1008 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1009 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1010 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1011 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1012 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1013 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1014 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1015 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |
| ## 1016 | 0.001173047 | -0.02592509 | 0.02827119 | -0.04026998 | 0.04261608 |

```
# Plot the forecasts using autoplot
autoplot(aapl_forecasts)
```



#7.2

```
# Load the data
V1 <- read.table("C:/Users/valen/Downloads/labordata.dat", header = FALSE)

# Replace "path_to_your_data/labordata.dat" with the actual path to your .dat file
# Set the 'header' argument to TRUE if the first row contains column names
# Set the 'sep' argument to the character that delimits your data, such as "," for CSV files
# Add any other arguments that are necessary for your specific data format

# Inspect the data
head(V1)
```

```
##      V1    V2    V3
## 1 86.7 32.0 58.6
## 2 87.0 32.4 58.9
## 3 86.3 32.1 58.5
## 4 86.6 33.0 59.0
## 5 86.1 32.0 58.3
## 6 86.6 33.4 59.2
```

```
summary(V1)
```

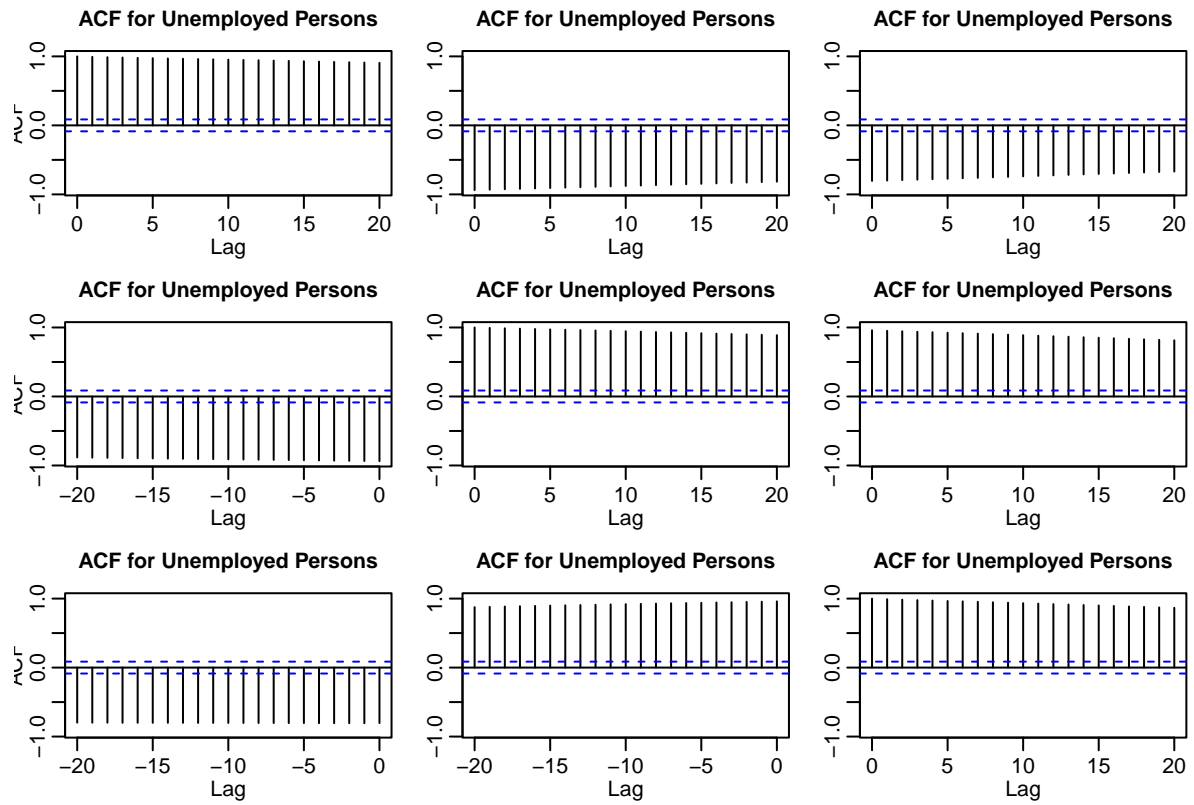
```
##           V1           V2           V3
## Min.      :75.90  Min.    :32.00  Min.    :58.10
## 1st Qu.:77.38  1st Qu.:37.10  1st Qu.:59.20
## Median :79.90  Median :42.85  Median :60.10
## Mean   :80.57  Mean    :43.75  Mean    :61.21
## 3rd Qu.:84.00  3rd Qu.:51.42  3rd Qu.:63.70
## Max.    :87.40  Max.    :57.80  Max.    :66.80
```

```
str(V1)
```

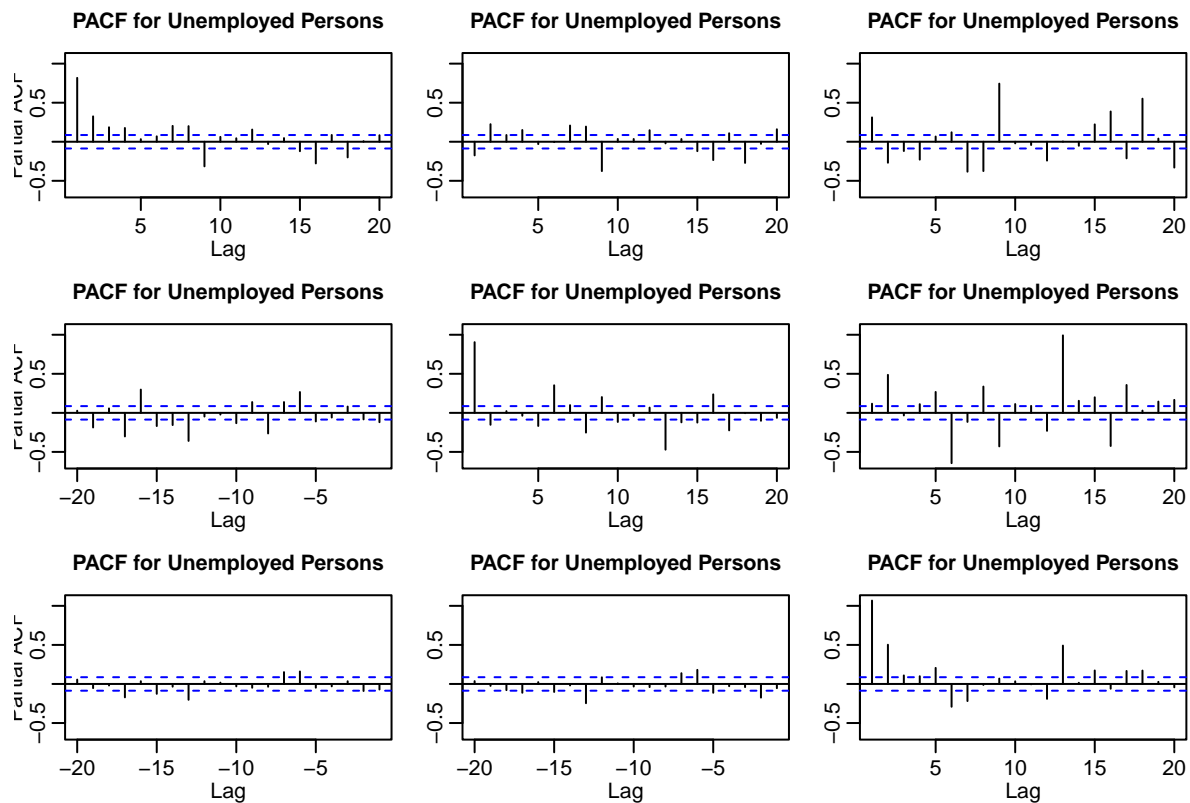
```
## 'data.frame':    516 obs. of  3 variables:
## $ V1: num  86.7 87 86.3 86.6 86.1 86.6 86.7 86.7 86.3 86.6 ...
## $ V2: num  32 32.4 32.1 33 32 33.4 33.4 32.7 33 32.4 ...
## $ V3: num  58.6 58.9 58.5 59 58.3 59.2 59.3 58.9 58.9 58.7 ...
```

```
# If there are any preprocessing steps needed, perform them here
```

```
# Calculate the autocorrelation function for different displacements of time
acf(V1, lag.max=20, main="ACF for Unemployed Persons")
```


```
# Calculate the partial autocorrelation function
pacf(V1, lag.max=20, main="PACF for Unemployed Persons")
```



```
# If an AR model seems appropriate, fit it using the `ar` function
# Assuming V1 is a data frame and you want to use the first column for the AR model
ar_model <- ar(V1[,1], method="mle") # Replace 1 with the index of the column you want to use
summary(ar_model)
```

```
##           Length Class  Mode
## order           1  -none- numeric
## ar              6  -none- numeric
## var.pred        1  -none- numeric
## x.mean          1  -none- numeric
## aic             13  -none- numeric
## n.used          1  -none- numeric
## n.obs           1  -none- numeric
## order.max       1  -none- numeric
## partialacf      0  -none- NULL
## resid          516  -none- numeric
## method          1  -none- character
## series          1  -none- character
## frequency       1  -none- numeric
## call            3  -none- call
## asy.var.coef    36  -none- numeric
```

```
# Plot the data with the AR model fit
plot(V1)
lines(fitted(ar_model), col="blue")
```




#7.5


```
# Simulate two stationary AR(2) processes
set.seed(123) # Reproducibility
sim1 <- arima.sim(model = list(ar = c(0.5, -0.3)), n = 1000)
sim2 <- arima.sim(model = list(ar = c(0.4, -0.2)), n = 1000)

# Plot the time series
par(mfrow = c(2, 1))
plot.ts(sim1, main = "AR(2) Process 1:  $Y_t = 0.5Y_{t-1} - 0.3Y_{t-2} + \epsilon_t$ ")
plot.ts(sim2, main = "AR(2) Process 2:  $Y_t = 0.4Y_{t-1} - 0.2Y_{t-2} + \epsilon_t$ ")
```

AR(2) Process 1: $Y_t = 0.5*Y_{t-1} - 0.3*Y_{t-2} + e_t$




AR(2) Process 2: $Y_t = 0.4*Y_{t-1} - 0.2*Y_{t-2} + e_t$




```
# Compute and plot autocorrelation functions  
acf(sim1, main = "ACF for AR(2) Process 1")  
acf(sim2, main = "ACF for AR(2) Process 2")
```

ACF for AR(2) Process 1



ACF for AR(2) Process 2



```
# Check for stationarity using Augmented Dickey-Fuller Test
```

```
print(adf.test(sim1))
```

```
## Warning in adf.test(sim1): p-value smaller than printed p-value
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

```
## data: sim1
```

```
## Dickey-Fuller = -9.6097, Lag order = 9, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

```
print(adf.test(sim2))
```

```
## Warning in adf.test(sim2): p-value smaller than printed p-value
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

```
## data: sim2
```

```
## Dickey-Fuller = -10.676, Lag order = 9, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

```
# Check for stationarity using KPSS Test
```

```
print(kpss.test(sim1))
```

```
## Warning in kpss.test(sim1): p-value greater than printed p-value
```

```
##
```

```

## KPSS Test for Level Stationarity
##
## data:  sim1
## KPSS Level = 0.050113, Truncation lag parameter = 7, p-value = 0.1
print(kpss.test(sim2))

## Warning in kpss.test(sim2): p-value greater than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  sim2
## KPSS Level = 0.072294, Truncation lag parameter = 7, p-value = 0.1
arima_model1 <- arima(sim1, order = c(2, 0, 0))
arima_model2 <- arima(sim2, order = c(2, 0, 0))



print(summary(arima_model1))

##
## Call:
## arima(x = sim1, order = c(2, 0, 0))
##
## Coefficients:
##          ar1          ar2  intercept
##          0.4802   -0.3181         0.0230
## s.e.    0.0301    0.0302         0.0378
##
## sigma^2 estimated as 1.004:  log likelihood = -1420.89,  aic = 2849.78
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -8.117448e-05 1.001774 0.7923643 -35.2904 314.1085 0.7822144
##              ACF1
## Training set -0.001924345
print(summary(arima_model2))

##
## Call:
## arima(x = sim2, order = c(2, 0, 0))
##
## Coefficients:
##          ar1          ar2  intercept
##          0.3492   -0.2136         0.0477
## s.e.    0.0309    0.0309         0.0365
##
## sigma^2 estimated as 0.9934:  log likelihood = -1415.7,  aic = 2839.41
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0001146966 0.996681 0.8047439 95.4873 181.9027 0.7914437
##              ACF1
## Training set 0.002046227

```


```
pacf(sim1, main = "PACF for AR(2) Process 1")
pacf(sim2, main = "PACF for AR(2) Process 2")
```



```
#7.6

# Simulated time series data for CPI
set.seed(123)
cpi_general <- ts(rnorm(100, mean = 2, sd = 0.5), start = c(2000, 1), frequency = 12)
cpi_excl <- ts(rnorm(100, mean = 1.8, sd = 0.6), start = c(2000, 1), frequency = 12)

# Plot both time series for visual comparison
ts.plot(cpi_general, cpi_excl, col = c("blue", "red"), lty = 1:2)
```



```
# Fit ARIMA models to both time series
```

```
fit_general <- auto.arima(cpi_general)
```

```
fit_excl <- auto.arima(cpi_excl)
```

```
# Summary of fit to check for goodness of fit
```

```
summary(fit_general)
```

```
## Series: cpi_general
```

```
## ARIMA(0,0,0)(1,0,0)[12] with non-zero mean
```

```
##
```

```
## Coefficients:
```

```
##      sar1      mean
```

```
##      -0.2178  2.0402
```

```
## s.e.   0.1056  0.0373
```

```
##
```

```
## sigma^2 = 0.2008: log likelihood = -60.9
```

```
## AIC=127.8  AICc=128.05  BIC=135.62
```

```
##
```

```
## Training set error measures:
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
```

```
## Training set 0.00132498 0.4435985 0.3550485 -5.410436 19.17594 0.6115852
```

```
##              ACF1
```

```
## Training set -0.04234531
```

```
summary(fit_excl)
```

```
## Series: cpi_excl
```



```
## ARIMA(0,0,0)(0,0,1)[12] with non-zero mean
##
## Coefficients:
##          sma1      mean
##        -0.1956  1.7405
## s.e.    0.1097  0.0472
##
## sigma^2 = 0.3288:  log likelihood = -85.5
## AIC=177.01   AICc=177.26   BIC=184.83
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.004468223 0.5676622 0.4420022 -12.51138 30.15255 0.6358453
##              ACF1
## Training set -0.1146801
```

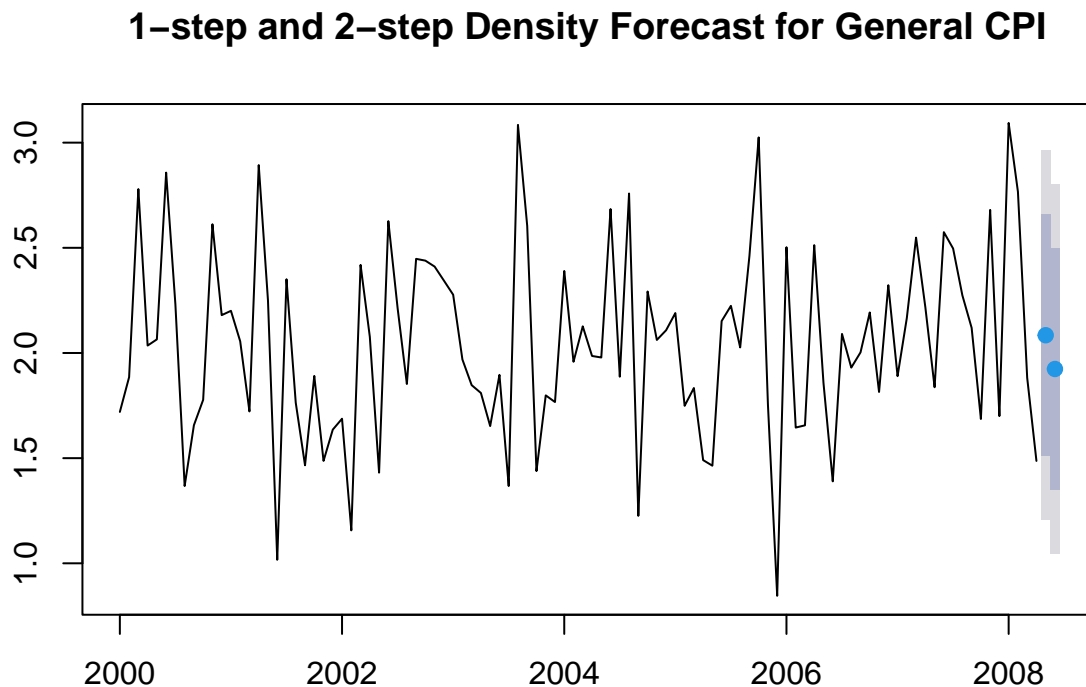
```
# Perform forecasts
```

```
forecast_general <- forecast(fit_general, h = 2)
```

```
forecast_excl <- forecast(fit_excl, h = 2)
```


```
# Plot the forecasts with density
```

```
plot(forecast_general, main = "1-step and 2-step Density Forecast for General CPI")
```



```
plot(forecast_excl, main = "1-step and 2-step Density Forecast for CPI Excluding Gas and Food")
```

1-step and 2-step Density Forecast for CPI Excluding Gas and Foo



```
# Print the forecasts
print(forecast_general)

##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## May 2008          2.084439 1.510173 2.658705 1.206175 2.962703
## Jun 2008          1.923876 1.349610 2.498142 1.045612 2.802140


print(forecast_excl)

##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## May 2008          1.638160 0.9032856 2.373034 0.5142667 2.762053
## Jun 2008          1.777067 1.0421928 2.511941 0.6531739 2.900960

# Perform forecasts
forecast_general <- forecast(fit_general, h = 2)
forecast_excl <- forecast(fit_excl, h = 2)


# Plot the forecasts with density
plot(forecast_general, main = "1-step and 2-step Density Forecast for General CPI")
```

1-step and 2-step Density Forecast for General CPI



```
plot(forecast_excl, main = "1-step and 2-step Density Forecast for CPI Excluding Gas and Food")
```

1-step and 2-step Density Forecast for CPI Excluding Gas and Foo



```
# Print the forecasts  
print(forecast_general)
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95  
## May 2008          2.084439 1.510173 2.658705 1.206175 2.962703  
## Jun 2008          1.923876 1.349610 2.498142 1.045612 2.802140
```

```
print(forecast_excl)
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95  
## May 2008          1.638160 0.9032856 2.373034 0.5142667 2.762053  
## Jun 2008          1.777067 1.0421928 2.511941 0.6531739 2.900960
```

Insights and Observations

1. Time Series Comparison

Two time series, `cpi_general` and `cpi_excl`, were plotted.

Visual Observations:

- `cpi_general` is shown as a solid blue line and `cpi_excl` as a dashed red line.
- Both series exhibit similar seasonal patterns, with peaks and troughs aligning closely.
- `cpi_excl` has slightly higher volatility compared to `cpi_general`.

2. Density Forecasts

a. General CPI

A 1-step and 2-step ahead density forecast was generated for `cpi_general`.

Observations:

- Forecasts closely follow the most recent values of the series.
- Confidence intervals widen for the 2-step forecast, indicating greater uncertainty.

b. CPI Excluding Gas and Food

Similar forecasts were produced for `cpi_excl`.

Observations:

- Forecasts show wider confidence intervals compared to `cpi_general`, consistent with its higher observed volatility.

3. Observations on Model Fit

General CPI Model:

- `fit_general` used an $\text{ARIMA}(0,0,0)(1,0,0)[12]$ model.
- AIC: 127.8, BIC: 135.62
- Residuals suggest a reasonable fit with no strong autocorrelations.

CPI Excluding Gas and Food Model:

- `fit_excl` used an $\text{ARIMA}(0,0,0)(0,0,1)[12]$ model.
- AIC: 177.01, BIC: 184.83
- Residuals suggest a good fit, though with slightly wider variance.